

Keeping torch.compile on (hooks as a custom op) recovers most of the speed the compile-off hook engine gave up:
steered generation lands within 2% of plain vLLM on the 27B / vLLM 0.27.1

