

Embedding-stream injection on DeepSeek-V4-Flash (hyper-connection architecture) costs no decode time: decode-step time of embedding replacement and embedding add vs no steering, one distinct vector per request
vLLM 0.27.1 · TP4 · B200 x4 · fp8 + fp4 experts · kv fp8_ds_mla · decode-step time = (wall at 80 new tokens – wall at 40) / 40, error bars = min-max over interleaved repeats · vllm-metamodel 1.1.0.post3

