

Decode runs at the no-steering speed with one distinct vector per request: on Qwen3.6-27B the decode-step time of norm-matched addition and embedding replacement is within $\pm 1.7\%$ of no steering, eager and with CUDA graphs
(on Qwen3-1.7B a 0.5–2 s call cannot resolve per-step differences: error bars = min-max over interleaved repeats; the CUDA-graph evidence there is the hook count, 3 vs 41 invocations)
decode-step time = (wall at 80 new tokens – wall at 40) / 40 · 27B: min of 2 repeats, 1.7B: paired median of 3 · 96-token prompt · bf16 · 1× B200 · vllm-metamodel 1.1.0.post2

