





Merging the LoRA into the weights removes the LoRA kernels from every decode step
Qwen3.6-27B: LoRA 115.3 ms -> merged 82.3 ms per decode step at B=512 (29% less, vLLM 0.19.0) · Qwen3-1.7B: LoRA 9.5 ms -> merged 6.2 ms per decode step at B=512 (34% less, vLLM 0.27.1)



LoRA-capable engine, no adapter adapter merged into the weights (LoRA-capable engine) adapter merged, plain engine
rank-64 LoRA on every request plain engine (enable_lora=False)

decode step = (wall for 40 new tokens - wall for 1 token) / 39; one steering vector per request; CUDA graphs (FULL_DECODE_ONLY); rank 64, rsLoRA alpha 16, ||sBA||/||W|| = 0.5% per module