

Prefix caching with per-request steering stays exact (steered blocks salted) and saves 7-15% on a dense-attention model; the hybrid 27B gains nothing because vLLM cannot cache its recurrent (GatedDeltaNet) state

