

In-engine readout with early exit reads layer 42 of 64 for 1,024 texts 2.5× faster than stock vllm-lens capture and 1.4× faster than the HF early-exit forward
Qwen3.6-27B bf16 · 1× B200 · 1,024 texts of 116 tokens on average · cosine of every position with a per-text direction, max over the last 5 · min of repeats

