

In-engine readout with early exit reads layer 18 of 28 for 1,024 texts 4.7× faster than stock vllm-lens capture and 1.5× faster than the HF early-exit forward
Qwen3-1.7B bf16 · 1× B200 · 1,024 texts of 116 tokens on average · cosine of every position with a per-text direction, max over the last 5 · min of repeats

