

Reading one layer's residual stream out of vLLM: stock capture cost grows with every request, the fork's readout stays at the no-hook ceiling  
prefill-only re-encode (max\_tokens=1) of 96-136-token texts · bf16 · 1× B200 · lines = min over repeats

