



Qwen/Qwen3.6-27B on 1x B200: one steering vector per request (RL-rollout style) runs 3-38x faster than stock vllm-lens 1.1.0 with the indexed hook, vectorised injection and CUDA graphs

96-token prompt, 40 new tokens per request, bf16, one steering vector per request (layer 1, one prompt position, distinct random unit vector per request). Ceilings = the same vLLM without any steering.



- stock vllm-lens 1.1.0 (per-layer key scan, eager forced)
- vllm-lens-metamodel: indexed hook, eager
- vllm-lens-metamodel: indexed + vectorised apply, eager
- vllm-lens-metamodel: indexed + vectorised + CUDA graphs