


Newer vLLM generates faster on the same B200, and the fork's per-request steering stays close to plain vLLM on every version
Qwen3.6-27B: plain vLLM 0.27.1 is 1.05x 0.19.0; fork+graphs at 83% of plain · Qwen3-1.7B: plain vLLM 0.28.0 is 1.47x 0.16.0; fork+graphs at 81% of plain



96-token prompt, 40 new tokens, one distinct steering vector per request (layer 1, marker token), best of 2 repeats; plain = VLLM_LENS_DISABLE=1 with vLLM's default torch.compile + CUDA graphs