

Hidden-state readout with early exit stays the cheapest way to read layer L on every vLLM version tested (1,024 texts of 96-136 tokens, layer 42/64 on the 27B and 18/28 on the 1.7B, CUDA-graph engine)

