

Per-request steering on Qwen3.6-27B (1× B200): vllm-metamodel generates as fast as plain vLLM; stock vllm-lens is 38× slower at batch 2,048
96-token prompt, 40 new tokens, bf16; lower is better

