

Per-request steering on Qwen3-1.7B (1× B200): vllm-metamodel matches plain vLLM; stock vllm-lens is 59× slower at batch 1,024
96-token prompt, 40 new tokens, bf16; lower is better

